

# Low-Resource English→Kannada News Translation: Step-by-Step Guide

This tutorial walks through building an English–Kannada news-domain MT system using synthetic data and efficient fine-tuning. We cover corpus creation with LLMs, dataset management on Hugging Face, preprocessing, PEFT fine-tuning on limited GPU (T4), evaluation, and sharing the final model/dataset. Code examples and tips are provided for reproducibility and memory efficiency.

## 1. Create a Synthetic English–Kannada Parallel Corpus

1. **Collect English news data.** Gather monolingual English news text from open sources (e.g. RSS feeds, Wikipedia news, or Kaggle news datasets). For example, you might scrape headlines or paragraphs using Python `requests` or use public datasets (news summaries, etc.) as your *source sentences*.
2. **Generate Kannada translations with an LLM or translation model.** Use a large language model (LLM) to translate English sentences into Kannada. You can use the OpenAI GPT API (e.g. ChatGPT/GPT-4) with a prompt like:

*"This is an English to Kannada translation. Provide only the Kannada sentence."*

Or use a free/open-source multilingual model. For example, Hugging Face's M2M100 (418M) supports Kannada. In Python:

```
from transformers import pipeline
translator = pipeline("translation", model="facebook/m2m100_418M",
src_lang="en", tgt_lang="kn")
en_sentences = ["The economy is growing steadily.",
"Heavy rains are expected tomorrow."] # example
kn_sentences = [translator(s, max_length=100)[0]['translation_text'] for s in
en_sentences]
```

Each English sentence maps to a Kannada translation. Ensure the model outputs only the translation (no extra text). Alternatively, if using GPT-3.5/4 via API, send batches of sentences with a clear instruction (no explanations). As research shows, LLM-generated synthetic data can significantly improve low-resource MT (even if noisy) <sup>1</sup> <sup>2</sup>.

1. **Save as CSV/JSON.** Structure the parallel data in a CSV or JSONL with fields like `english` and `kannada`. For example:

```
import pandas as pd
data = {'english': en_sentences, 'kannada': kn_sentences}
```

```
df = pd.DataFrame(data)
df.to_csv("en_kn_news.csv", index=False)
```

*Tip:* If you have a large list of sentences, process in batches. Include a header row (e.g. `english,kannada`). You can also save as JSON lines: each line `{"english": "...", "kannada": "..."}.`

## 2. Upload the Dataset to Hugging Face

1. **Create a HF account and dataset repo.** On [huggingface.co](https://huggingface.co), sign up/log in. Click your profile → **New Dataset**, give it a name (e.g. `username/en-kn-news-synthetic`), and choose public or private.

The screenshot shows the Hugging Face web interface for creating a new dataset repository. The header includes the Hugging Face logo, a search bar, and navigation links for Models, Datasets, Spaces, Docs, Solutions, and Pricing. A user profile dropdown menu is open, showing options like 'Profile', 'mishig', 'New Model', 'New Dataset', 'New Space', 'New Organization', 'New AutoNLP Project', 'Inference Dashboard', 'Settings', and 'Sign Out'. The main form is titled 'Create a new dataset repository' and includes a subtitle 'A repository contains all dataset files, including the revision history.' The form has fields for 'Owner' (mishig), 'Dataset name' (New dataset name), and 'License'. There are two radio button options for visibility: 'Public' (selected) and 'Private'. The 'Public' option description says 'Anyone on the internet can see this dataset. Only you (personal dataset) or members of your organization (organization dataset) can commit.' The 'Private' option description says 'Only you (personal dataset) or members of your organization (organization dataset) can see and commit to this dataset.' A 'Create dataset' button is at the bottom.

Figure: Creating a new dataset repository on Hugging Face.

1. **Push data via UI or Python.** In the new repo's **Files and versions** tab, click *Add file* → *Upload file*, and upload your CSV (you may ZIP it if large). Alternatively, use Python:

```
from huggingface_hub import login
from datasets import load_dataset

login(token="YOUR_HF_TOKEN")      # find/create on your HF account settings
dataset = load_dataset("csv", data_files="en_kn_news.csv")['train']
dataset.push_to_hub("yourusername/en-kn-news-synthetic")
```

This logs in and pushes the dataset (you must have created the repo first). The `push_to_hub()` function makes the dataset available on HF (see code example 3).

## 3. Preprocess and Format for Training

1. **Load with Datasets.** In your training notebook (e.g. Colab), install and import:

```
!pip install transformers datasets sentencepiece
from datasets import load_dataset
raw_data = load_dataset("yourusername/en-kn-news-synthetic", split="train")
```

1. **Clean and split.** Optionally filter out empty/malformed entries. Then split into train/validation:

```
raw_data = raw_data.filter(lambda x: x['english'].strip() != "" and
x['kannada'].strip() != "")
splits = raw_data.train_test_split(test_size=0.1, seed=42)
train_ds, val_ds = splits['train'], splits['test']
```

1. **Tokenization.** Choose a tokenizer for your model. For a bilingual seq2seq model (e.g. M2M100), use its `AutoTokenizer`. For a decoder-only model (Llama/Qwen), you may prepend a prompt. Example (using a T5-style tokenizer):

```
from transformers import AutoTokenizer
tokenizer = AutoTokenizer.from_pretrained("facebook/m2m100_418M")
def preprocess(examples):
    inputs = examples['english']
    targets = examples['kannada']
    model_inputs = tokenizer(inputs, truncation=True, padding="longest")
    with tokenizer.as_target_tokenizer():
        labels = tokenizer(targets, truncation=True, padding="longest")
    model_inputs["labels"] = labels["input_ids"]
    return model_inputs
train_tok = train_ds.map(preprocess, batched=True)
val_tok = val_ds.map(preprocess, batched=True)
```

If using a causal model (e.g. LLaMA), you might concatenate like `"Translate English to Kannada: \n<English>"` as input and use the generated part as target. Adjust `tokenizer` accordingly. The goal is to produce `input_ids` and `labels` for training.

## 4. Fine-Tune a Small Model with LoRA on Colab T4

1. **Setup environment in Colab.** Ensure GPU (T4) selected. Install libraries:

```
!pip install transformers datasets peft accelerate bitsandbytes
```

1. **Load base model** (1–2B params). For example, Meta’s LLaMA 1B HF version, using 8-bit to fit GPU memory:

```
from transformers import AutoModelForCausalLM, AutoTokenizer
model_name = "mesolitica/llama-1b-hf-32768-fpf" # LLaMA 1B HuggingFace repo
tokenizer = AutoTokenizer.from_pretrained(model_name)
model = AutoModelForCausalLM.from_pretrained(
    model_name,
    load_in_8bit=True, # or load_in_4bit if available
```

```

        device_map="auto"
    )

```

*Tip:* Use `load_in_8bit=True` (via [bitsandbytes](#)) to reduce GPU memory. Also set `device_map="auto"` to shard model across devices (Colab T4 has 16GB).

1. **Apply LoRA (PEFT).** Configure LoRA to keep the base model frozen and only train small adapters

4 :

```

from peft import LoraConfig, get_peft_model
lora_config = LoraConfig(
    task_type="CAUSAL_LM",
    inference_mode=False,
    r=8, # rank
    lora_alpha=32,
    lora_dropout=0.05,
    target_modules=["q_proj", "v_proj"] # adjust to model architecture
)
model = get_peft_model(model, lora_config)

```

This wraps the model to only train low-rank updates; ~99% of parameters stay frozen 4 .

1. **Prepare Trainer.** Use `transformers.Trainer` with gradient accumulation to save memory:

```

from transformers import Trainer, TrainingArguments

training_args = TrainingArguments(
    output_dir="en-kn-lora",
    per_device_train_batch_size=4, # reduce if OOM
    per_device_eval_batch_size=4,
    gradient_accumulation_steps=4, # simulate larger batch
    learning_rate=3e-4,
    fp16=True, # mixed precision
    max_steps=5000, # or number of epochs
    logging_steps=100,
    evaluation_strategy="steps",
    eval_steps=500,
    save_total_limit=2
)
trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=train_tok,
    eval_dataset=val_tok,
    data_collator=lambda x: {'input_ids': torch.tensor([d['input_ids'] for d
in x]),
                            'labels': torch.tensor([d['labels'] for d in
x])}
)

```

```
)
trainer.train()
```

**Memory tips:** T4 GPUs have ~16GB VRAM. Use `fp16=True` and small `per_device_train_batch_size` (e.g. 1-4). Increase `gradient_accumulation_steps` to simulate a larger effective batch. Optionally enable `model.gradient_checkpointing = True` or use `load_in_4bit` if supported.

1. **Save the fine-tuned model.** After training, save the LoRA adapter:

```
model.save_pretrained("en_kn_lora_adapter")
```

To share it later, you can push to Hugging Face hub (see Step 8).

## 5. Evaluate Translation Quality

1. **Generate translations.** Compare the base model vs. the LoRA-fine-tuned model on a held-out test set. Example generation loop:

```
from transformers import AutoModelForCausalLM

# Reload base and fine-tuned model
base_model = AutoModelForCausalLM.from_pretrained(model_name,
load_in_8bit=True, device_map="auto")
tuned_model = AutoModelForCausalLM.from_pretrained("en_kn_lora_adapter",
load_in_8bit=True, device_map="auto")

test_sentences = val_ds['english'][:100] # use part of validation as test
def translate(model, texts):
    result = []
    for text in texts:
        inputs = tokenizer(text, return_tensors="pt").to(model.device)
        out = model.generate(**inputs, max_new_tokens=128)
        result.append(tokenizer.decode(out[0], skip_special_tokens=True))
    return result

preds_base = translate(base_model, test_sentences)
preds_tuned = translate(tuned_model, test_sentences)
refs = val_ds['kannada'][:100]
```

1. **Compute BLEU/ChrF.** Use Hugging Face's `evaluate` library:

```
import evaluate
bleu = evaluate.load("sacrebleu")
chrF = evaluate.load("chrF")

base_bleu = bleu.compute(predictions=preds_base, references=[[r] for r in
refs])
```

```
tuned_bleu = bleu.compute(predictions=preds_tuned, references=[[r] for r in
refs])
base_chrf = chrf.compute(predictions=preds_base, references=refs)
tuned_chrf = chrf.compute(predictions=preds_tuned, references=refs)
print(f"Base BLEU: {base_bleu['score']:.1f}, Tuned BLEU:
{tuned_bleu['score']:.1f}")
print(f"Base ChrF: {base_chrf['score']:.1f}, Tuned ChrF:
{tuned_chrf['score']:.1f}")
```

These metrics quantify translation quality. As expected, the LoRA-tuned model should outperform the base model (BLEU/ChrF increases). For example, you might observe something like Base BLEU  $\approx 10$ , Tuned BLEU  $\approx 16$ , Base ChrF  $\approx 25$ , Tuned ChrF  $\approx 32$  (illustrated below).

【24†】 (Figure: Example metric improvements for English→Kannada translation after LoRA fine-tuning.)

1. **Example translations.** Inspect some outputs:

```
for en, ref, base_pred, tuned_pred in zip(test_sentences[:5], refs[:5],
preds_base[:5], preds_tuned[:5]):
    print("EN:", en)
    print("REF:", ref)
    print("BASE:", base_pred)
    print("TUNED:", tuned_pred)
    print("-"*40)
```

Document cases where LoRA tuning fixes errors (reordering, terminology) compared to the base model.

## 6. Visualize and Document Improvements

Present your evaluation results clearly:

- **Plot metrics.** As above, a bar chart can visualize Base vs. Tuned BLEU/ChrF. (See embedded example.) This highlights gains.
- **Translation examples.** Include a table or list of English inputs, reference Kannada, and both model outputs. Emphasize improvements (e.g. better fluency or correct named entities) after tuning.
- **Training logs.** Show loss/accuracy curves from Trainer logs if available. Comment on convergence.

Ensure each figure/table is captioned. The example bar chart above shows **Base** vs **Finetuned** on BLEU/ChrF, illustrating the boost from synthetic data+LoRA.

## 7. Scripts and Colab Tips

- **Full Colab script.** Combine the above steps into a Colab notebook or Python scripts. Break cells by step (data prep, tokenization, training). Use `%pip install` at top. Ensure `!pip install` commands and imports are included.

- **Memory saving.** Besides `fp16` and 8-bit loading, also try `--max_length 128` to limit sequence length. Use `dataset = dataset.map(tokenizer, batched=True)` to speed up tokenization. If encountering OOM, drop batch size or accumulate more.
- **Checkpointing.** Save LoRA adapter periodically (Trainer's `save_steps`) so you don't lose progress on disconnection.
- **Reproducibility.** Set random seeds (`TrainingArguments.seed`) and document versions of libraries.

## 8. Documentation and Sharing

1. **Write a clear README.** On your HF dataset/model repos, include a README with dataset/model description, usage examples, licenses, and links. For example, say "English news → Kannada news translations, synthetically generated with GPT/M2M100; fine-tuned on LLaMA-1B with LoRA."
2. **Tagging.** Use HF's tagging tool or manually add tags in the card (languages, tasks: translation, news domain, synthetic data, etc.) to improve discoverability.
3. **Upload the model.** Create a new model repo on Hugging Face (e.g. `yourusername/en-kn-news-lora`). Then in Python or CLI, push the adapter:

```
huggingface-cli login
```

```
from huggingface_hub import HfApi
api = HfApi()
api.upload_folder(
    folder_path="en_kn_lora_adapter",
    path_in_repo="",
    repo_id="yourusername/en-kn-news-lora",
    token="YOUR_HF_TOKEN"
)
```

Alternatively, use `model.push_to_hub("yourusername/en-kn-news-lora")` if the model is a `transformers.PreTrainedModel` (ensure `model.save_pretrained()` called first). Include `tokenizer` files in the repo.

1. **Citing your work.** In your repository's README (and any submission), cite relevant references. For example, note that LoRA is used to reduce trainable parameters <sup>4</sup>, and synthetic data from LLMs boosts low-resource MT <sup>1</sup>.

By following these steps and including code, you will have a reproducible pipeline: from synthetic data generation to a LoRA-tuned Kannada translation model, all sharable on Hugging Face. This demonstrates both technical skills (dataset creation, efficient fine-tuning) and practical considerations (memory limits, documentation) suitable for a candidate assignment.

**References:** Hugging Face documentation for datasets and PEFT (LoRA) <sup>3</sup> <sup>4</sup>; research on LLMs for low-resource MT <sup>1</sup> <sup>2</sup>.

1 2 **Scaling Low-Resource MT via Synthetic Data Generation with LLMs**

<https://arxiv.org/html/2505.14423v1>

3 **Upload a dataset to the Hub — datasets 1.16.0 documentation**

[https://huggingface.co/docs/datasets/v1.16.0/upload\\_dataset.html](https://huggingface.co/docs/datasets/v1.16.0/upload_dataset.html)

4 **LoRA**

[https://huggingface.co/docs/peft/main/en/conceptual\\_guides/lora](https://huggingface.co/docs/peft/main/en/conceptual_guides/lora)